

- 1) Let v be an n -size vector of integers. Show that after a pre-processing in $O(n)$ time, we can answer to the following type of queries $q(i,j)$ in $O(1)$ -time: is the subvector $v[i..j]$ sorted?

Solution 1: For every index k , let $inv[k]$ denote the number of indices $s < k$ such that $v[s] > v[s+1]$. It can be constructed in $O(n)$ by scanning the vector v once. Now, to answer to $q(i,j)$, we output YES if either $i=j$ or $inv[j] - inv[i] = 0$.

Solution 2: We partition the vector in maximal sorted sub-vectors. It can be done in $O(n)$ by scanning once vector v . Now, to answer to $q(i,j)$, we just need to check in $O(1)$ whether i,j are in the same range of this partition.

Comment: we could also check whether $v[i..j]$ is either sorted by non-increasing values or sorted by non-decreasing values.

- 2) We are given as input an n -size integer vector v . We are allowed to pre-process this vector so as to answer, as fast as possible, to the following type of queries $q(i,j)$: “what is the number of elements $v[p]$, p between i and j , such that $v[p]$ is an even number?”.

- a. Show that, if v is fixed, then we can achieve $O(n)$ pre-processing time and $O(1)$ query time.

We use a classic ‘partial sum’ trick. Let u be the n -size boolean vector so that $u[i] = 1$ if and only if $v[i]$ is even. In an auxiliary vector w , we store in $w[i]$ the number of even elements $v[q]$, for q between 0 and i . Since $w[0] = u[0]$ and $w[i] = w[i-1] + u[i]$, then we can compute the vector w in $O(n)$ time by dynamic programming. Now, for answering a query $q(i,j)$, it suffices to output $w[j] - w[i] + u[i]$.

- b. We now allow the vector v to be dynamically modified: at any time step, one can change the value stored in an arbitrary position i of the vector. Show that we can achieve $O(n)$ pre-processing time and $O(\sqrt{n})$ query time.

We use Mo’s trick. Specifically, we partition the vector in \sqrt{n} contiguous blocks of size \sqrt{n} each. In a separate \sqrt{n} -size vector s , we store the number of even elements in each block. Note that, upon modifying the value of $v[i]$, if i lies in the j th block then we only need to modify $s[j]$ (we decrement this value by 1 if the former value of $v[i]$ was an even number, then we increment this value by 1 if the new value of $v[i]$ is even), that only takes $O(1)$ time. Now, in order to answer to a query $q(i,j)$, we do as follows: let $B_a, B_{a+1}, \dots, B_{a+t}$ the blocks fully between i and j (we can find these blocks simply by looking at the blocks to which i and j belong, respectively). We start summing all values $s[a+p]$, for p between 0 and t . Since t is at most \sqrt{n} , the latter can be done in $O(\sqrt{n})$ time. Then, we increment the result by 1 for each even element $v[q]$, $q \geq i$, of B_{a-1} and we end up incrementing the result also by 1 for each even element $v[q]$, $j \geq q$, of B_{a+t+1} . Since each block has at most \sqrt{n} elements, it also takes $O(\sqrt{n})$ time.

- 3) Presentation of Merge Sort (not seen yet in class).

- 4) We are given as input an n -size integer vector v .

- a. Show that we can compute in $O(n \log(n))$ the number of inversions, that is, the number of pairs (r,s) such that $r < s$ and $v[r] > v[s]$.

It suffices to adapt Merge sort. We cut the vector in two halves and we apply recursively our algorithms on both halves so that: a) we counted the number of inversions in both halves; and b) we sorted both halves. Then, during a classical merge of both halves in

one sorted vector (interclasare), we can count the number of inversions with one element in each half. For that, consider the i th element of the left half. Let it be put in position $i+j$ in the final sorted vector. Then, j elements on the right half were smaller than it. Therefore, we count j inversions with the right half for this element.

- b. Show that after pre-processing the vector in $O(n \cdot \log(n))$, we can answer to the following type of queries $q(i)$ in $O(1)$: "what is the number of inversions between i and $n-1$, that is, the number of pairs (r,s) such that $i \leq r < s$ and $v[r] > v[s]$?"

We apply the algorithm for counting the number of inversions in a vector (see the previous question). It runs in $O(n \cdot \log(n))$ time. Doing so, we actually computed something stronger: namely, we computed for every index i the number $inv[i]$ of indices $j > i$ such that $v[i] > v[j]$. Now, we apply a classic "partial sum trick", namely: let $u[i] = \sum_{k=0 \dots i} inv[k]$. In order to answer to a query $q(i)$, it suffices to output $u[n-1] - u[i] + inv[i]$. The pre-processing time is in $O(n \cdot \log(n))$, while the query time is in $O(1)$.

- c. We are allowed to pre-process this vector so as to answer, as fast as possible, to the following type of queries $q(i,j)$: "what is the number of inversions between i and j , that is, the number of pairs (r,s) such that $i \leq r < s \leq j$ and $v[r] > v[s]$?"

To solve our range query problem, we can now combine the above algorithm with Mo's trick. Specifically, we partition the vector in \sqrt{n} contiguous blocks of size \sqrt{n} each.

i) We create a first $\sqrt{n} \times \sqrt{n}$ matrix M_0 , so that: if $i \leq j$, then $M_0[i,j]$ is the number of pairs (r,s) such that $r < s$ is in block i , s is in block j , and $v[r] > v[s]$. If $i=j$, then we can compute $M_0[i,i]$ in $O(\sqrt{n} \cdot \log(n))$ time by using our above algorithm. If $i < j$, then we count in $O(\sqrt{n} \cdot \log(n))$ time the number $inv(i,j)$ of inversions in the vector obtained from the concatenation of the i th and j th block; then, $M_0[i,j] = inv(i,j) - M_0[i,i] - M_0[j,j]$ (we could also easily adapt our previous algorithm to this case). The total runtime is in $O(n \cdot \sqrt{n} \cdot \log(n))$.

ii) Then, we use a classic partial sum trick. Specifically, for $i \leq j$, let $M_1[i,j]$ be the sum of all values $M_0[i,j']$, $i \leq j' \leq j$. We can compute the matrix M_1 in $O(n)$ time by dynamic programming.

iii) Now, we create another $n \times \sqrt{n}$ matrix M_2 so that $M_2[i,j]$ is the number of elements smaller than $v[i]$ in the j th block. Being given a sorted copy of B_j we can compute $M_2[i,j]$ in $O(\log(n))$, simply by doing a binary search for $v[i]$. Therefore, the total runtime is in $O(n \cdot \sqrt{n} \cdot \log(n))$.

iv) Finally, we create the $n \times \sqrt{n}$ matrix M_3 so that $M_3[i,j]$ is the sum of all values $M_2[i,j']$, $0 \leq j' \leq j$. We can compute the matrix M_3 in $O(n \cdot \sqrt{n})$ time by dynamic programming.

In order to answer to a query $q(i,j)$, we do as follows: let $B_a, B_{a+1}, \dots, B_{a+t}$ the blocks fully between i and j (we can find these blocks simply by looking at the blocks to which i and j belong, respectively). We start summing all values $M_1[a+p, a+t]$, for p between 0 and t . Since t is at most \sqrt{n} , the latter can be done in $O(\sqrt{n})$ time. For each element $v[q]$, $q \geq i$, of B_{a-1} we increment the output by $M_3[q, a+t] - M_3[q, a-1]$. Similarly, for each element $v[q]$, $j \geq q$, of B_{a+t+1} , we increment the output by $M_3[q, a+t] - M_3[q, a-1]$. We are left computing the number of inversions in $[i,j] \cap B_{a-1} \cup B_{a+t+1}$ that can be done in $O(\sqrt{n} \cdot \log(n))$ time.