

- 1) A word is made of open '(' and closed ')' parentheses. A subword is made of consecutive letters. Prove that we can compute a longest well formed subword in $O(n)$ time.

The key observation is as follows. Let us assume that we store a counter c , initially set to 0, which we increment/decrement after each open/closed parenthesis. Let i be the first index such that c becomes negative (if it exists). Then, a well-formed subword is fully between positions 0 and $i-1$ or it is fully between positions $i+1$ and $n-1$ (i.e., there is no overlap). In particular, we may restart the algorithm each time the counter c becomes negative, and then we never need to go back in the vector.

More precisely, the algorithm works as follows. We maintain a starting position s (initially, $s = 0$), an optimal length L (initially, $L=0$) and a counter c (initially set to 0). We increment/decrement c after each open/closed parenthesis. Let i denote our current position.

* If $c=0$, then we found a new well-formed subword: $L = \max\{L, i-s+1\}$

* If $c < 0$, then we reset $c=0$, $s=i+1$

- 2) Let $v[]$ be a vector of n positive integers. Being given an integer k , compute in $O(n)$ time the minimum length of a subvector $v[i..j]$ such that the sum of all its elements is at least k .

For every index i , let $\min[i]$ be the least index j such that $v[i..j]$ satisfies the desired property (the sum of all its elements is at least k). Since every element is positive, we have: $\min[i+1] \geq \min[i]$. Therefore, we can compute the vector $\min[]$ as follows:

* we use a partial sum trick: for every index i , let $s[i]$ be equal to $v[0] + v[1] + \dots + v[i]$. We can compute vector $s[]$ in $O(n)$ time by dynamic programming.

* set $i = j = 0$ (here, i represents the first index for which we do not know $\min[i]$; j is a lower bound for $\min[i]$).

```
* while i < n do:
    while j < n and s[j]-s[i]+v[i] < k do:
        j=j+1
    min[i] = j
    i = i+1
```

* finally, we iterate over all indices i such that $\min[i] < n$ and we output the minimum of $\min[i] - i + 1$.

- 3) A frequency-stack is a data structure supporting the following operations:
 - * `empty()`: asserts whether the structure is empty
 - * `push(e)`: add a new element e to the structure (possibly the repetition of a previous element)
 - * `pop()`: returns and deletes a most frequent element in the structure. All occurrences of e are removed at once.

Describe an implementation such that every operation can be performed in expected $O(1)$ time.

We store a "list of list" F . More precisely, F is a doubly-linked list, whose every element is a pair $(r, L(r))$. The list $L(r)$ contains every element repeated r times in the structure. Furthermore, F is ordered by increasing number of repetitions, i.e.: $F = [r_1, r_2, \dots, r_q]$ with $r_1 < r_2 < \dots < r_q$.

*`empty()`: we just check whether F is empty.

*`push(e)`: There are several cases and subcases

a) Case e was never added to the structure before

Let r_1 be the head of F

If $r_1=1$ then we add e to $L(1)$. Otherwise, we prepend 1 in F , and we set $L(1) = [e]$

b) Case e was already inserted r times to the structure before.

We remove e from $L(r)$.

If r is the tail of F , or the successor r' of r in F is not $r+1$, then we insert $r+1$ in F immediately after r , and we set $L(r+1) = [e]$. Otherwise, we add e to $L(r+1)$.

*`pop()`: Let r_q be the tail of F . We remove any element from $L(r_q)$. If $L(r_q)$ becomes empty, then we remove

rq from F.

Every operation can be done in $O(1)$ time if, for any element e , we can:

- compute in $O(1)$ time the number of repetitions of e in the structure (or decide that e was never added before)

- if there are r repetitions of e , access in $O(1)$ time to the position of e in list $L(r)$.

We can store this information in an auxiliary Hash-table. Thus, any operation can be done in expected $O(1)$ time.

- 4) Let $u[]$ and $v[]$ be vectors of length n . We say that u and v are anagrams if they have the exact same elements (counted with multiplicities), but not necessarily in the same positions. Example: $[1,1,2,3,1]$ and $[1,2,1,3,1]$ are anagrams. Propose an algorithm in expected $O(n)$ time in order to decide whether u, v are anagrams.

We create a Hash-table H whose keys are the elements of vector u , and such that $H[e]$ is the number of repetitions of e in u . This can be done in expected $O(n)$ time simply by scanning vector u once. Then, we consider each element of v sequentially, from $j=0$ to $j=n-1$. Let $e = v[j]$ be the current element. If e is not a key of H , then we stop (e is in v but not in u , therefore u and v are not anagrams). Otherwise, we decrement $H[e]$ and, if $H[e]$ drops to 0, we remove e from H . If we end scanning v without stopping, then we accept.

- 5) Let $v[]$ be a vector of n positive integers.
- a) Compute the maximum length of a subvector $v[i..j]$ such that all its elements are pairwise distinct.

We scan vector v from left to right. Let $e = v[j]$ be the current element. During the scan, we store in an auxiliary Hash-table all elements e' encountered on previous positions $0, 1, \dots, j-1$ and the last position where each element was found. In particular, we have access to $H[e]$, the largest index $j' < j$ such that $v[j'] = v[j] = e$, with the convention that $H[e] = -1$ if no such index exists. During the scan, we compute a vector $b[]$ such that, for every j , $b[j]$ equals the least index i such that $v[i..j]$ satisfies the desired property (all its elements are distinct). Note that computing $b[]$ is actually sufficient in order to solve our problem. We compute all values $b[j]$ by induction:

- * if $j=0$, then $b[j] = 0$

- * otherwise, $b[j] = \max\{b[j-1], H[e]+1\}$, where $e = v[j]$.

b) Compute the maximum length of a subvector $v[i..j]$ that satisfies Fibonacci recurrence, i.e.: for every k between $i+2$ and j , we must have $v[k] = v[k-1] + v[k-2]$.

First we create an auxiliary vector $\text{fibo}[]$ of length $n-2$ such that: $\text{fibo}[k] = 1$ if and only if $v[k+2] = v[k] + v[k+1]$ (else, $\text{fibo}[k] = 0$). This can be done in $O(n)$ time by scanning the vector v once.

Then, we compute $\text{max-fibo}[]$, also of length $n-2$, such that $\text{max-fibo}[i]$ equals the maximum length of a Fibonacci subvector that starts in position i .

- * If $\text{fibo}[i] = 0$, then $\text{max-fibo}[i] = 2$.

- * If $i=n-3$ and $\text{fibo}[i]=1$, then $\text{max-fibo}[i] = 3$

- * if $i < n-3$ and $\text{fibo}[i] = 1$, then $\text{max-fibo}[i] = 1 + \text{max-fibo}[i+1]$

The total running time is in $O(n)$.

Remark: The runtime for 5b) is slightly better than for 5a), for which the running time was also $O(n)$ but only in expectation. The reason for that is the Fibonacci property is local, whereas checking that all elements are pairwise distinct is a global property. For global properties, it is often difficult to avoid using hash-table, whereas simpler tricks often work for local properties.

- 6) Let $v[]$ be a vector of n positive integers. Compute, in expected $O(n)$ time, the smallest positive integer e that is NOT an element of v .

It suffices to consider integers e between 1 and $n+1$ (since we only have n elements of e , the smallest missing integer must be in this range). For that, we store all elements of v in a hash-table H . Then, for e from 1 to $n+1$, we search for e in H .